

EXPERIMENT- 3

AIM: Write a program to implement Artificial Neural Network for MNIST dataset.

CODE and OUTPUT:

```
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.metrics import confusion_matrix

import keras
from keras.datasets import mnist
from keras.layers import Dense
from keras.models import Sequential
from matplotlib import pyplot as plt
from random import randint

# Preparing the dataset
# Setup train and test splits
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# Making a copy before flattening for the next code-segment which displays images
x_train_drawing = x_train
print("X_Train:",x_train[0])
print("y_train:",y_train[0])

print("X_Train Shape:",x_train.shape)
print("y_train Shape:",y_train.shape)

image_size = 784 # 28 x 28
x_train = x_train.reshape(x_train.shape[0], image_size)
x_test = x_test.reshape(x_test.shape[0], image_size)

print("After reshaping")
print("X_Train Shape:",x_train.shape)
print("x_test Shape:",x_test.shape)


# Convert class vectors to binary class matrices
num_classes = 10
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

After reshaping
X_Train Shape: (60000, 784)
x_test Shape: (10000, 784)

```
print(y_train.shape)
print(y_train[0])
```

```
(60000, 10)
[0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

```
for i in range(64):
    ax = plt.subplot(8, 8, i+1)
    ax.axis('off')
    plt.imshow(x_train_drawing[randint(0, x_train.shape[0])], cmap='Greys')
```



```
model = Sequential()

# The input layer requires the special input_shape parameter which should match
# the shape of our training data.
model.add(Dense(units=32, activation='sigmoid', input_shape=(image_size,)))
model.add(Dense(units=num_classes, activation='softmax'))
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 32)	25120
dense_1 (Dense)	(None, 10)	330
Total params: 25,450		
Trainable params: 25,450		
Non-trainable params: 0		

```
model.compile(optimizer="sgd", loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(x_train, y_train, batch_size=128, epochs=100, verbose=True, validation_split=.1)
```

```
Epoch 1/100
422/422 [=====] - 1s 3ms/step - loss: 0.1584 - accuracy: 0.9538 - val_loss: 0.1714 - val_accuracy: 0.9517
Epoch 2/100
422/422 [=====] - 1s 2ms/step - loss: 0.1527 - accuracy: 0.9551 - val_loss: 0.1730 - val_accuracy: 0.9508
Epoch 3/100
422/422 [=====] - 1s 2ms/step - loss: 0.1607 - accuracy: 0.9525 - val_loss: 0.1670 - val_accuracy: 0.9517
...
Epoch 97/100
422/422 [=====] - 1s 2ms/step - loss: 0.1289 - accuracy: 0.9612 - val_loss: 0.1630 - val_accuracy: 0.9550
Epoch 98/100
422/422 [=====] - 1s 2ms/step - loss: 0.1299 - accuracy: 0.9604 - val_loss: 0.1583 - val_accuracy: 0.9532
Epoch 99/100
422/422 [=====] - 1s 2ms/step - loss: 0.1292 - accuracy: 0.9606 - val_loss: 0.1627 - val_accuracy: 0.9495
Epoch 100/100
422/422 [=====] - 1s 2ms/step - loss: 0.1284 - accuracy: 0.9609 - val_loss: 0.1578 - val_accuracy: 0.9545
```

```
loss,accuracy = model.evaluate(x_test, y_test, verbose=True)
```

```
313/313 [=====] - 0s 955us/step - loss: 0.1911 - accuracy: 0.9430
```

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['training', 'validation'], loc='best')
plt.show()
```

